

In [5]: '''Trains a simple deep NN on the MNIST dataset.

Gets to 98.40% test accuracy after 20 epochs
(there is *a lot* of margin for parameter tuning).
2 seconds per epoch on a K520 GPU.
'''

```
from tensorflow import keras
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop

batch_size = 128
num_classes = 10
epochs = 20

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

60000 train samples
10000 test samples
Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401920
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 512)	262656
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 10)	5130

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

Epoch 1/20
469/469 [=====] - 6s 11ms/step - loss: 0.2469 - accuracy: 0.9240 - val_loss: 0.1025 - val_accuracy: 0.9670
Epoch 2/20
469/469 [=====] - 6s 12ms/step - loss: 0.1014 - accuracy: 0.9696 - val_loss: 0.0863 - val_accuracy: 0.9753
Epoch 3/20
469/469 [=====] - 6s 13ms/step - loss: 0.0744 - accuracy: 0.9772 - val_loss: 0.0698 - val_accuracy: 0.9804
Epoch 4/20
469/469 [=====] - 6s 12ms/step - loss: 0.0594 - accuracy: 0.9815 - val_loss: 0.0757 - val_accuracy: 0.9793
Epoch 5/20
469/469 [=====] - 6s 12ms/step - loss: 0.0510 - accuracy: 0.9853 - val_loss: 0.0722 - val_accuracy: 0.9802
Epoch 6/20
469/469 [=====] - 6s 12ms/step - loss: 0.0433 - accuracy: 0.9879 - val_loss: 0.0720 - val_accuracy: 0.9826
Epoch 7/20
469/469 [=====] - 6s 12ms/step - loss: 0.0394 - accuracy: 0.9883 - val_loss: 0.0827 - val_accuracy: 0.9813
Epoch 8/20
469/469 [=====] - 5s 12ms/step - loss: 0.0347 - accuracy: 0.9897 - val_loss: 0.0850 - val_accuracy: 0.9821
Epoch 9/20
469/469 [=====] - 6s 13ms/step - loss: 0.0329 - accuracy: 0.9905 - val_loss: 0.0926 - val_accuracy: 0.9819
Epoch 10/20
469/469 [=====] - 6s 13ms/step - loss: 0.0283 - accuracy: 0.9918 - val_loss: 0.0886 - val_accuracy: 0.9832
Epoch 11/20
469/469 [=====] - 6s 12ms/step - loss: 0.0280 - accuracy: 0.9918 - val_loss: 0.0906 - val_accuracy: 0.9825
Epoch 12/20
469/469 [=====] - 6s 12ms/step - loss: 0.0270 - accuracy: 0.9926 - val_loss: 0.0861 - val_accuracy: 0.9845
Epoch 13/20
469/469 [=====] - 6s 12ms/step - loss: 0.0232 - accuracy: 0.9932 - val_loss: 0.1013 - val_accuracy: 0.9835

```

Epoch 14/20
469/469 [=====] - 6s 13ms/step - loss: 0.0242 - accuracy: 0.
9935 - val_loss: 0.1047 - val_accuracy: 0.9831
Epoch 15/20
469/469 [=====] - 6s 13ms/step - loss: 0.0217 - accuracy: 0.
9941 - val_loss: 0.1067 - val_accuracy: 0.9842
Epoch 16/20
469/469 [=====] - 6s 13ms/step - loss: 0.0220 - accuracy: 0.
9939 - val_loss: 0.1230 - val_accuracy: 0.9816
Epoch 17/20
469/469 [=====] - 6s 13ms/step - loss: 0.0199 - accuracy: 0.
9946 - val_loss: 0.1218 - val_accuracy: 0.9820
Epoch 18/20
469/469 [=====] - 6s 12ms/step - loss: 0.0191 - accuracy: 0.
9949 - val_loss: 0.1161 - val_accuracy: 0.9835
Epoch 19/20
469/469 [=====] - 6s 12ms/step - loss: 0.0181 - accuracy: 0.
9954 - val_loss: 0.1294 - val_accuracy: 0.9816
Epoch 20/20
469/469 [=====] - 6s 13ms/step - loss: 0.0189 - accuracy: 0.
9950 - val_loss: 0.1289 - val_accuracy: 0.9827
Test loss: 0.12885063886642456
Test accuracy: 0.982699990272522

```

In []:

```

---
title: Assignment 1
subtitle: Computer performance, reliability, and scalability calculation
author: Shaquiel Pashtunyar
03/19/2023
---

## 1.2

#### a. Data Sizes

| Data Item | Size per Item | |
|---|---|---|
| 128 character message. | 128 Byte |
| 1024x768 PNG image | 4.72 MB | https://toolstud.io/photo |
| 1024x768 RAW image | 1.18 MB |
| HD (1080p) HEVC Video (15 minutes) | 2500MB | https://toolstud.io/video |
| HD (1080p) Uncompressed Video (15 minutes) | 15600 MB |
| 4K UHD HEVC Video (15 minutes) | 2550 MB | 170MB per 60 seconds (1 |
| 4k UHD Uncompressed Video (15 minutes) | 3660 MB |
| Human Genome (Uncompressed) | 6.27 GB | https://www.google.com/s |

#### b. Scaling

| | Size | # HD | |
|---|---|---|---|
| Daily Twitter Tweets (Uncompressed) | 64GB | 1 | 128*500 64,00bytes=0.0 |
| Daily Twitter Tweets (Snappy Compressed) | 37.5 | 1 | 64/1.7 |
| Daily Instagram Photos | 354TB | 36 | 4.72*75000000 |
| Daily YouTube Videos | 500TB | 50 | 500 * 60 = 30000 mir |
| Yearly Twitter Tweets (Uncompressed) | 23.36TB | 3 | 64 GB * 365 = 23360 GB |
| Yearly Twitter Tweets (Snappy Compressed) | 13.7 | 2 | 37.65 GB * 365 = 13742 |
| Yearly Instagram Photos | 41062.5TB | 4107 | 112.5 TB (daily) * 365 |
| Yearly YouTube Videos | 182500TB | 18250 | 500 TB * 365 = 182500 |

#### c. Reliability

```

	# HD	# Failures
Twitter Tweets (Uncompressed)	3	0.0255
Twitter Tweets (Snappy Compressed)	2	0.017
Instagram Photos	4107	~35
YouTube Videos	18250	~3155

Failure rate = 0.85% (<https://www.backblaze.com/b2/hard-drive-test-data.html>)

$0.0085 * 3 = 0.0255$

$0.0085 * 2 = 0.017$

$0.0085 * 4107 = 34.9095$

$0.0085 * 18250 = 155.125$

d. Latency

	One Way Latency
Los Angeles to Amsterdam	140 ms
Low Earth Orbit Satellite	600 ms
Geostationary Satellite	240 ms
Earth to the Moon	2560 ms
Earth to Mars	13 minutes

<https://wondernetwork.com/pings>

<https://www.omniaccess.com/leo/#:~:text=The%20GE0%20Latency%20is%20of,and%20an%20esser>

<https://www.satsig.net/latency.htm>

https://en.wikipedia.org/wiki/Earth%E2%80%93Moon%E2%80%93Earth_communication#:~:text=F

<https://blogs.esa.int/mex/2012/08/05/time-delay-between-mars-and-earth>

```
In [7]: #
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements. See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#

import sys
from random import random
from operator import add

from pyspark.sql import SparkSession

if __name__ == "__main__":
    """
    Usage: pi [partitions]
    """
```

```

spark = SparkSession\
    .builder\
    .appName("PythonPi")\
    .getOrCreate()

partitions = int(sys.argv[1]) if len(sys.argv) > 1 else 2
n = 100000 * partitions

def f(_):
    x = random() * 2 - 1
    y = random() * 2 - 1
    return 1 if x ** 2 + y ** 2 <= 1 else 0

count = spark.sparkContext.parallelize(range(1, n + 1), partitions).map(f).reduce(
print("Pi is roughly %f" % (4.0 * count / n))

spark.stop()

```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[7], line 34
    26 """
    27     Usage: pi [partitions]
    28 """
    29 spark = SparkSession\
    30     .builder\
    31     .appName("PythonPi")\
    32     .getOrCreate()
--> 34 partitions = int(sys.argv[1]) if len(sys.argv) > 1 else 2
    35 n = 100000 * partitions
    37 def f(_):

ValueError: invalid literal for int() with base 10: '-f'

```

In []:

In []: