

```
In [2]: ###Assignment 6
        #DSC650 4/23/2023
        #Shaquiel Pashtunyar
```

Assignment 6

Using section 5.1 in Deep Learning with Python as a guide (listing 5.3 in particular), create a ConvNet model that classifies images in the MNIST digit dataset. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

```
In [5]: #Importing
```

```
In [49]: from keras import layers
        from keras import models
        from keras.datasets import mnist
        from keras.utils import to_categorical
        import numpy as np
        import matplotlib.pyplot as plt
        import pandas as pd
        import os
        import pandas as pd
```

```
In [12]: #Mnist imported
```

```
In [13]: # calling a model
        model = models.Sequential()
        model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(28, 28, 1)))
        model.add(layers.MaxPooling2D((2, 2)))
        model.add(layers.Conv2D(64, (3,3), activation='relu'))
        model.add(layers.MaxPooling2D((2,2)))
        model.add(layers.Conv2D(64, (3,3), activation='relu'))

        # add a classifier on top of Convnet
        model.add(layers.Flatten())
        model.add(layers.Dense(64, activation='relu'))
        model.add(layers.Dense(10, activation = 'softmax'))

        # view summary
        model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling 2D)	(None, 13, 13, 32)	0
conv2d_4 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_3 (MaxPooling 2D)	(None, 5, 5, 64)	0
conv2d_5 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_2 (Dense)	(None, 64)	36928
dense_3 (Dense)	(None, 10)	650

=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
=====

```
In [14]: #Splitting the data into training and target
```

```
In [15]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

```
In [16]: # shuffle training set
for _ in range(5):
    indexes = np.random.permutation(len(train_images))

    train_images = train_images[indexes]
    train_labels = train_labels[indexes]

    # put 10,000 aside for validation
    val_images = train_images[:10000,:]
    val_labels = train_labels[:10000,:]

    # keep the rest in training set
    train_images2 = train_images[10000,:,:]
    train_labels2 = train_labels[10000,:,:]

    # view their shape
    train_images2.shape, val_images.shape
```

```
Out[16]: ((50000, 28, 28, 1), (10000, 28, 28, 1))
```

```
In [18]: #model compiling
```

```
model.compile(optimizer='rmsprop',
               loss='categorical_crossentropy',
               metrics=['accuracy'])

history = model.fit(train_images2, train_labels2, epochs=5, batch_size=64,
                    validation_data = (val_images, val_labels))
```

Epoch 1/5

782/782 [=====] - 18s 22ms/step - loss: 0.1986 - accuracy: 0.9387 - val_loss: 0.1153 - val_accuracy: 0.9655

Epoch 2/5

782/782 [=====] - 17s 22ms/step - loss: 0.0510 - accuracy: 0.9840 - val_loss: 0.0584 - val_accuracy: 0.9830

Epoch 3/5

782/782 [=====] - 20s 26ms/step - loss: 0.0356 - accuracy: 0.9891 - val_loss: 0.0485 - val_accuracy: 0.9852

Epoch 4/5

782/782 [=====] - 19s 24ms/step - loss: 0.0268 - accuracy: 0.9920 - val_loss: 0.0352 - val_accuracy: 0.9888

Epoch 5/5

782/782 [=====] - 21s 26ms/step - loss: 0.0209 - accuracy: 0.9935 - val_loss: 0.0405 - val_accuracy: 0.9877

```
In [20]: #file path needed for download
```

```
os.getcwd()
```

```
Out[20]: 'C:\\Users\\spashtunyar\\Documents\\School\\dsc650\\dsc650\\assignments\\assignment06'
```

```
In [37]: train_loss = history.history['loss']
         val_loss = history.history['val_loss']
```

```
epochs = range(1, len(history.history['loss']) + 1)
```

```
plt.plot(epochs, train_loss, 'bo', label = 'Training Loss')
```

```
plt.plot(epochs, val_loss, 'b', label = 'Validation Loss')
```

```
plt.title('Training & Validation Losses')
```


```
plt.xlabel('Epochs')
```

```
plt.ylabel('Loss')
```

```
plt.legend()
```

```
plt.savefig(r'C:\\Users\\spashtunyar\\Documents\\School\\dsc650\\dsc650\\assignments\\assignment06\\loss.png')
```


```
plt.show()
```



```
In [36]: #Accuracy
train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

epochs = range(1, len(history.history['accuracy']) + 1)

plt.plot(epochs, train_acc, 'bo', label = 'Training Accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation Accuracy')
plt.title('Training & Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.savefig(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment6\figs\fig6_10.png')
plt.show()
```



```
In [39]: # retrain model & evaluate for 3 epochs
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs = 3, batch_size = 64)
results = model.evaluate(test_images, test_labels)
```

Epoch 1/3
938/938 [=====] - 25s 25ms/step - loss: 0.0116 - accuracy: 0.9965
Epoch 2/3
938/938 [=====] - 25s 26ms/step - loss: 0.0092 - accuracy: 0.9973
Epoch 3/3
938/938 [=====] - 25s 26ms/step - loss: 0.0076 - accuracy: 0.9980
313/313 [=====] - 2s 5ms/step - loss: 0.0435 - accuracy: 0.9911

```
In [40]: results
```

```
Out[40]: [0.043479882180690765, 0.991100013256073]
```

```
In [42]: #Loss and accuracy
history.history
```

```
Out[42]: {'loss': [0.011624673381447792, 0.009243286214768887, 0.007589162793010473],
          'accuracy': [0.9964666962623596, 0.9973499774932861, 0.9979833364486694]}
```

```
In [43]: model.save(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment
```

```
In [44]: prediction_results = model.predict(test_images)

313/313 [=====] - 2s 5ms/step
```

```
In [45]: prediction_results
```

```
Out[45]: array([[8.4818215e-13, 1.4013637e-12, 7.9418064e-13, ..., 1.0000000e+00,
        7.1448720e-14, 8.8554679e-13],
        [1.0670370e-12, 3.0961589e-13, 1.0000000e+00, ..., 6.2001581e-20,
        6.3564592e-16, 4.9383507e-21],
        [3.2855095e-15, 1.0000000e+00, 1.3975719e-14, ..., 2.2907494e-10,
        2.7650401e-12, 3.6749812e-15],
        ...,
        [1.0209669e-26, 6.3787823e-16, 4.9082052e-22, ..., 2.7280451e-16,
        3.5719909e-18, 1.6070612e-15],
        [2.5390151e-14, 3.2301263e-19, 7.8710421e-22, ..., 3.7945682e-18,
        2.5481697e-07, 1.1942707e-15],
        [2.1948858e-13, 5.9157579e-20, 4.1198930e-18, ..., 8.3996679e-24,
        2.9717780e-16, 5.2695157e-21]], dtype=float32)
```

```
In [46]: with open(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment
f.write('Training Loss: {}'.format(str(history.history['loss'])))
f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
f.write('\nTest Loss: {}'.format(results[0]))
f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
In [47]: preds = pd.DataFrame(prediction_results,
        columns = ['0','1','2','3','4','5','6','7','8','9'])

preds.to_csv(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignm
```

Assignment part 2

```
In [48]: #importing cifar10
from keras.datasets import cifar10
```

```
In [50]: #Splitting data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 8s 0us/step

```
In [51]: # preprocess data
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

# put 10,000 aside for validation
x_val = x_train[-10000:]
y_val = y_train[-10000:]
x_train = x_train[:-10000]
y_train = y_train[:-10000]
```

```
In [52]: # Create the model
model = models.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation = 'softmax'))

# view summary
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_6 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_4 (MaxPooling 2D)	(None, 15, 15, 32)	0
conv2d_7 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_5 (MaxPooling 2D)	(None, 6, 6, 64)	0
conv2d_8 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_6 (MaxPooling 2D)	(None, 2, 2, 64)	0
flatten_2 (Flatten)	(None, 256)	0
dense_4 (Dense)	(None, 64)	16448
dense_5 (Dense)	(None, 10)	650
=====		
Total params: 73,418		
Trainable params: 73,418		
Non-trainable params: 0		

```
In [54]: model.compile(optimizer='rmsprop',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])


history = model.fit(x_train, y_train, epochs=100,
                  validation_data = (x_val, y_val), verbose=0)
```

```
In [55]: train_loss = history.history['loss']
val_loss = history.history['val_loss']


epochs = range(1, len(history.history['loss']) + 1)

plt.plot(epochs, train_loss, 'bo', label = 'Training Loss')
```

```
plt.plot(epochs, val_loss, 'b', label = 'Validation Loss')
plt.title('Training & Validation Losses')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.savefig(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment6\figs\fig6_10.png')
plt.show()
```



```
In [56]: train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(1, len(history.history['accuracy']) + 1)
plt.plot(epochs, train_acc, 'bo', label = 'Training Accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation Accuracy')
plt.title('Training & Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.savefig(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment6\figs\fig6_11.png')
plt.show()
```

```
In [57]: # retrain model & evaluate
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# preprocess data
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs=10)
results = model.evaluate(x_test, y_test)
```

```

Epoch 1/10
1563/1563 [=====] - 39s 24ms/step - loss: 0.8312 - accuracy:
0.7800
Epoch 2/10
1563/1563 [=====] - 37s 24ms/step - loss: 0.7477 - accuracy:
0.7891
Epoch 3/10
1563/1563 [=====] - 40s 26ms/step - loss: 0.7026 - accuracy:
0.7936
Epoch 4/10
1563/1563 [=====] - 47s 30ms/step - loss: 0.6806 - accuracy:
0.7970
Epoch 5/10
1563/1563 [=====] - 42s 27ms/step - loss: 0.6500 - accuracy:
0.8025
Epoch 6/10
1563/1563 [=====] - 44s 28ms/step - loss: 0.6398 - accuracy:
0.8050
Epoch 7/10
1563/1563 [=====] - 41s 26ms/step - loss: 0.6412 - accuracy:
0.8010
Epoch 8/10
1563/1563 [=====] - 39s 25ms/step - loss: 0.6225 - accuracy:
0.8078
Epoch 9/10
1563/1563 [=====] - 39s 25ms/step - loss: 0.6123 - accuracy:
0.8088
Epoch 10/10
1563/1563 [=====] - 37s 24ms/step - loss: 0.6165 - accuracy:
0.8118
313/313 [=====] - 3s 8ms/step - loss: 1.2173 - accuracy: 0.6
453

```

```
In [59]: model.save(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignmer
```

```
In [60]: prediction_results = model.predict(x_test)
```

```
313/313 [=====] - 3s 8ms/step
```

```
In [61]: prediction_results
```

```
Out[61]: array([[4.21460310e-04, 1.52559573e-04, 2.89890799e-03, ...,
1.02881808e-04, 2.96914834e-03, 1.25233841e-04],
[6.28627837e-04, 3.70398047e-04, 3.56287654e-16, ...,
1.48146096e-14, 9.99000967e-01, 1.30031585e-08],
[5.94272576e-02, 1.90348569e-02, 2.61096051e-04, ...,
8.35009923e-05, 9.08882439e-01, 8.88891425e-03],
...,
[3.65121878e-07, 4.80077063e-08, 3.01973009e-03, ...,
1.10382631e-01, 4.64879292e-07, 1.73881290e-05],
[9.73920669e-06, 3.36405193e-03, 1.74554007e-03, ...,
7.68269529e-04, 3.40715123e-09, 3.31380056e-06],
[4.94112008e-11, 1.59739350e-18, 7.89580952e-08, ...,
9.99655366e-01, 5.83787319e-19, 4.79023188e-18]], dtype=float32)
```

```
In [62]: #Writing results in text file
with open(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment
f.write('Training Loss: {}'.format(str(history.history['loss'])))
f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
```

```
f.write('\nTest Loss: {}'.format(results[0]))
f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
In [63]: preds = pd.DataFrame(prediction_results,
                             columns = ['0','1','2','3','4','5','6','7','8','9'])

preds.to_csv(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignm
```

```
In [1]: #Part 2 B
```

```
In [2]: from keras.datasets import cifar10
from keras.utils import to_categorical
from keras.preprocessing.image import ImageDataGenerator
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
In [5]: x_train.shape, y_train.shape
```

```
Out[5]: ((50000, 32, 32, 3), (50000, 1))
```

```
In [6]: # preprocess data
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

# put 10,000 aside for validation
x_val = x_train[-10000:]
y_val = y_train[-10000:]
x_train2 = x_train[:-10000]
y_train2 = y_train[:-10000]
```

```
In [7]: train_datagen = ImageDataGenerator(rescale=1./255,
                                           rotation_range=40,
                                           width_shift_range=0.2,
                                           height_shift_range=0.2,
                                           shear_range=0.2,
                                           zoom_range=0.2,
                                           horizontal_flip=True)

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow(x_train2, y_train2, batch_size=32)

validation_generator = train_datagen.flow(x_val, y_val, batch_size=32)
```

```
In [8]: from keras import models
from keras import layers
```

```
In [9]: # instantiate the model
# add dropout layer
model = models.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
```

```

model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation = 'softmax'))

# view summary
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
dense (Dense)	(None, 64)	16448
dense_1 (Dense)	(None, 10)	650
=====		
Total params: 73,418		
Trainable params: 73,418		
Non-trainable params: 0		

In [11]: `from keras import optimizers`

In [12]: `model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
loss='categorical_crossentropy',
metrics=['accuracy'])`

C:\Users\spashtunyar\Anaconda3\lib\site-packages\keras\optimizers\optimizer_v2\rmsprop.py:140: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
super().__init__(name, **kwargs)

In [13]: `history = model.fit_generator(train_generator,
steps_per_epoch=len(x_train2) / 32,
epochs = 30,
validation_data=validation_generator,
validation_steps=len(x_val) / 32)`

```
C:\Users\spashtunyar\AppData\Local\Temp\ipykernel_17548\7325378.py:1: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.  
    history = model.fit_generator(train_generator,
```

Epoch 1/30
1250/1250 [=====] - 120s 94ms/step - loss: 2.1608 - accuracy: 0.1845 - val_loss: 2.0044 - val_accuracy: 0.2605

Epoch 2/30
1250/1250 [=====] - 93s 74ms/step - loss: 1.9819 - accuracy: 0.2599 - val_loss: 1.9174 - val_accuracy: 0.2897

Epoch 3/30
1250/1250 [=====] - 68s 54ms/step - loss: 1.9184 - accuracy: 0.2846 - val_loss: 1.8656 - val_accuracy: 0.3120

Epoch 4/30
1250/1250 [=====] - 65s 52ms/step - loss: 1.8616 - accuracy: 0.3105 - val_loss: 1.8000 - val_accuracy: 0.3407

Epoch 5/30
1250/1250 [=====] - 57s 45ms/step - loss: 1.8219 - accuracy: 0.3250 - val_loss: 1.7632 - val_accuracy: 0.3502

Epoch 6/30
1250/1250 [=====] - 52s 42ms/step - loss: 1.7886 - accuracy: 0.3350 - val_loss: 1.7230 - val_accuracy: 0.3763

Epoch 7/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.7531 - accuracy: 0.3530 - val_loss: 1.6938 - val_accuracy: 0.3790

Epoch 8/30
1250/1250 [=====] - 54s 43ms/step - loss: 1.7335 - accuracy: 0.3632 - val_loss: 1.6988 - val_accuracy: 0.3836

Epoch 9/30
1250/1250 [=====] - 52s 41ms/step - loss: 1.7138 - accuracy: 0.3702 - val_loss: 1.6602 - val_accuracy: 0.4011

Epoch 10/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.6899 - accuracy: 0.3823 - val_loss: 1.6678 - val_accuracy: 0.3956

Epoch 11/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.6726 - accuracy: 0.3854 - val_loss: 1.6334 - val_accuracy: 0.4088

Epoch 12/30
1250/1250 [=====] - 51s 40ms/step - loss: 1.6576 - accuracy: 0.3959 - val_loss: 1.6215 - val_accuracy: 0.4082

Epoch 13/30
1250/1250 [=====] - 51s 40ms/step - loss: 1.6406 - accuracy: 0.3982 - val_loss: 1.5945 - val_accuracy: 0.4242

Epoch 14/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.6312 - accuracy: 0.4073 - val_loss: 1.5811 - val_accuracy: 0.4315

Epoch 15/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.6145 - accuracy: 0.4147 - val_loss: 1.5450 - val_accuracy: 0.4426

Epoch 16/30
1250/1250 [=====] - 51s 40ms/step - loss: 1.6048 - accuracy: 0.4162 - val_loss: 1.5401 - val_accuracy: 0.4425

Epoch 17/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.5919 - accuracy: 0.4244 - val_loss: 1.5398 - val_accuracy: 0.4463

Epoch 18/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.5847 - accuracy: 0.4245 - val_loss: 1.5134 - val_accuracy: 0.4678

Epoch 19/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.5737 - accuracy: 0.4310 - val_loss: 1.5199 - val_accuracy: 0.4492

Epoch 20/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.5555 - accuracy: 0.4404 - val_loss: 1.5258 - val_accuracy: 0.4523

```

Epoch 21/30
1250/1250 [=====] - 52s 42ms/step - loss: 1.5516 - accuracy:
0.4408 - val_loss: 1.4739 - val_accuracy: 0.4740
Epoch 22/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.5387 - accuracy:
0.4451 - val_loss: 1.4708 - val_accuracy: 0.4822
Epoch 23/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.5315 - accuracy:
0.4492 - val_loss: 1.4761 - val_accuracy: 0.4722
Epoch 24/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.5206 - accuracy:
0.4529 - val_loss: 1.4786 - val_accuracy: 0.4655
Epoch 25/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.5118 - accuracy:
0.4539 - val_loss: 1.4450 - val_accuracy: 0.4864
Epoch 26/30
1250/1250 [=====] - 52s 41ms/step - loss: 1.5075 - accuracy:
0.4549 - val_loss: 1.4584 - val_accuracy: 0.4829
Epoch 27/30
1250/1250 [=====] - 52s 41ms/step - loss: 1.5004 - accuracy:
0.4611 - val_loss: 1.4404 - val_accuracy: 0.4853
Epoch 28/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.4937 - accuracy:
0.4640 - val_loss: 1.4226 - val_accuracy: 0.4954
Epoch 29/30
1250/1250 [=====] - 50s 40ms/step - loss: 1.4840 - accuracy:
0.4648 - val_loss: 1.4246 - val_accuracy: 0.4959
Epoch 30/30
1250/1250 [=====] - 51s 41ms/step - loss: 1.4802 - accuracy:
0.4709 - val_loss: 1.4152 - val_accuracy: 0.4915

```


```

In [16]: train_loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(history.history['loss']) + 1)

plt.plot(epochs, train_loss, 'bo', label = 'Training Loss')
plt.plot(epochs, val_loss, 'b', label = 'Validation Loss')
plt.title('Training & Validation Losses')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.savefig(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment6\plots\training_validation_losses.png')
plt.show()


```



```
In [17]: train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

epochs = range(1, len(history.history['accuracy']) + 1)

plt.plot(epochs, train_acc, 'bo', label = 'Training Accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation Accuracy')
plt.title('Training & Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.savefig(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment6\figs\fig1.png')
plt.show()
```

```
In [18]: # retrain model & evaluate
train_generator = train_datagen.flow(x_train, y_train, batch_size=32)

model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit_generator(train_generator,
                             steps_per_epoch=len(x_train) / 32,
                             epochs = 16)

results = model.evaluate(x_test, y_test)
```

Epoch 1/16

C:\Users\spashtunyar\AppData\Local\Temp\ipykernel_17548\740027706.py:8: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

```
history = model.fit_generator(train_generator,
```

```

1562/1562 [=====] - 40s 25ms/step - loss: 1.4771 - accuracy:
0.4714
Epoch 2/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4708 - accuracy:
0.4739
Epoch 3/16
1562/1562 [=====] - 47s 30ms/step - loss: 1.4596 - accuracy:
0.4778
Epoch 4/16
1562/1562 [=====] - 47s 30ms/step - loss: 1.4533 - accuracy:
0.4809
Epoch 5/16
1562/1562 [=====] - 48s 31ms/step - loss: 1.4405 - accuracy:
0.4825
Epoch 6/16
1562/1562 [=====] - 51s 33ms/step - loss: 1.4372 - accuracy:
0.4871
Epoch 7/16
1562/1562 [=====] - 54s 34ms/step - loss: 1.4341 - accuracy:
0.4861
Epoch 8/16
1562/1562 [=====] - 54s 35ms/step - loss: 1.4223 - accuracy:
0.4940
Epoch 9/16
1562/1562 [=====] - 55s 35ms/step - loss: 1.4236 - accuracy:
0.4905
Epoch 10/16
1562/1562 [=====] - 55s 35ms/step - loss: 1.4146 - accuracy:
0.4976
Epoch 11/16
1562/1562 [=====] - 55s 35ms/step - loss: 1.4116 - accuracy:
0.4969
Epoch 12/16
1562/1562 [=====] - 57s 36ms/step - loss: 1.4000 - accuracy:
0.4995
Epoch 13/16
1562/1562 [=====] - 55s 35ms/step - loss: 1.4008 - accuracy:
0.5003
Epoch 14/16
1562/1562 [=====] - 59s 38ms/step - loss: 1.3937 - accuracy:
0.5051
Epoch 15/16
1562/1562 [=====] - 88s 56ms/step - loss: 1.3886 - accuracy:
0.5054
Epoch 16/16
1562/1562 [=====] - 61s 39ms/step - loss: 1.3816 - accuracy:
0.5086
313/313 [=====] - 3s 9ms/step - loss: 216.5171 - accuracy:
0.3698

```

```
In [19]: model.save(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignmer
```

```
In [20]: prediction_results = model.predict(x_test)
```

```
313/313 [=====] - 3s 9ms/step
```

```
In [21]: prediction_results
```

```
Out[21]: array([[9.0042595e-13, 0.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
          1.0000000e+00, 0.0000000e+00],
          [0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
          0.0000000e+00, 0.0000000e+00],
          [0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
          0.0000000e+00, 1.8758386e-21],
          ...,
          [0.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
          0.0000000e+00, 0.0000000e+00],
          [1.2673316e-34, 1.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
          0.0000000e+00, 0.0000000e+00],
          [0.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
          0.0000000e+00, 0.0000000e+00]], dtype=float32)
```

```
In [22]: with open(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignment
          f.write('Training Loss: {}'.format(str(history.history['loss'])))
          f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
          f.write('\nTest Loss: {}'.format(results[0]))
          f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
In [23]: preds = pd.DataFrame(prediction_results,
                               columns = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'])

preds.to_csv(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignm
```

Part 3

```
In [24]: !pip install opencv-python
```

```
Collecting opencv-python
  Downloading opencv_python-4.7.0.72-cp37-abi3-win_amd64.whl (38.2 MB)
    ----- 38.2/38.2 MB 23.4 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.17.3 in c:\users\spashtunyar\anaconda3\lib\si
te-packages (from opencv-python) (1.23.5)
Installing collected packages: opencv-python
Successfully installed opencv-python-4.7.0.72
```

```
In [47]: from tensorflow.keras.applications.resnet50 import ResNet50
          from tensorflow.keras.preprocessing import image
          from tensorflow.keras.applications.resnet50 import preprocess_input, decode_prediction
          import numpy as np
          import os, cv2
          import matplotlib.pyplot as plt
```

```
In [48]: # Load Model
          model = ResNet50(weights = 'imagenet')
```

```
In [49]: img_path = r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\assignmer
```

```
In [50]: images = os.listdir(img_path)
```

```
In [51]: # Get the filename for each image
          for i, name in enumerate(images):
              print(name)
```

6-1-metrics.txt
6-1-model.h5
6-1-predicitons.csv
6-2A-Accuracy.png
6-2A-Loss.png
6-2A-metrics.txt
6-2A-model.h5
6-2A-predicitons.csv
6-2B-Accuracy.png
6-2B-Loss.png
6-2B-metrics.txt
6-2B-predicitons.csv
6-2B_model.h5
accuracyPNG.png
assignment06PNG.png
image.txt

```
In [54]: # Write the predictions into a file
for i,name in enumerate(images):

    if name != '.ipynb_checkpoints':
        img = cv2.imread(img_path + '/' + name)
        img = cv2.resize(img, (224,224))
        x = image.img_to_array(img)
        x = np.expand_dims(x, axis=0)
        x = preprocess_input(x)

        preds = model.predict(x)

        decpr = name, decode_predictions(preds, top=3)[0]

        print(decpr)

        with open(r'C:\Users\spashtunyar\Documents\School\dsc650\dsc650\assignments\as
            f.write(decpr[0])
    else:
        pass
```

```
-----
error                                     Traceback (most recent call last)
Cell In[54], line 6
      4 if name != '.ipynb_checkpoints':
      5     img = cv2.imread(img_path + '/' + name)
----> 6     img = cv2.resize(img, (224,224))
      7     x = image.img_to_array(img)
      8     x = np.expand_dims(x, axis=0)

error: OpenCV(4.7.0) D:\a\opencv-python\opencv-python\opencv\modules\imgproc\src\resi
ze.cpp:4062: error: (-215:Assertion failed) !ssize.empty() in function 'cv::resize'
```

```
In [55]: #unsure here
```

```
In [ ]:
```